

# All-or-Nothing Atomicity

---

- An action is atomic
  - if there is no way for a higher layer to discover the **internal structure** of its implementation
- <> From the point of view of a procedure that invokes an atomic action
  - The atomic action always appears either to **complete** as anticipated, or to do **nothing**
  - atomic actions useful in **recovering** from **failures**

# Before-or-After Atomicity

---

- <> From the point of view of a concurrent thread
  - An atomic action acts as though it occurs either completely before or completely after every other concurrent atomic action
  - This consequence: makes atomic actions useful for coordinating concurrent threads
- Atomicity hides
  - not just the details of which steps form the atomic action
  - But the very fact that it has structure

# **All-or-Nothing and Before-or-After Atomicity**

---

- 1. Data abstraction
  - Hide the internal structure of data
- 2. Client/server organization
  - Hide the internal structure of major subsystems
- 3. Atomicity
  - Hide the internal structure of an action
- Enforce industrial-strength modularity
  - Guarantee absence of unanticipated interactions among components of a complex system
- The implementer's point of view
  - Painfully knowing the detail

# Atomic actions' **benevolent** side effects

---

- audit log
  - atomic actions that run into trouble record
    - the nature of the detected failure and
    - the recovery sequence
  - for later analysis
- Data management system when insert a record
  - rearrange the file into a better physical order
- Cache
- Garbage collection
- They are all hidden from upper levels

# Overall system fault tolerance model

---

- error-free operation不会出错的(不存在的):
  - All work goes according to expectations
  - The user initiates actions and the system confirms the actions by displaying messages to the user
- tolerated error（可容忍错误）：
  - The user who has initiated an action notices that the **system failed** before it confirmed completion of the action
  - when the system is operating again, **checks to see** whether or not it actually performed that action

# Overall system fault tolerance model

---

- untolerated error:
  - The system fails without the user noticing
  - the user does not realize that he or she should check or retry an action that the system may not have completed

# Disk Storage System Fault Tolerance Model

---

- A perfect-disk assumption
  - a disk never decays and that it has no hard errors
  - only one thing can go wrong: a system crash at just the wrong time

# Disk Storage System Fault Tolerance Model

---

- The fault tolerance model
  - error-free operation:
    - CAREFUL\_GET returns the result of the most recent call to CAREFUL\_PUT
    - at sector\_number on track, with status = OK.
  - detectable error:
    - The operating system crashes during a CAREFUL\_PUT
    - and corrupts the disk buffer in volatile storage
    - and CAREFUL\_PUT writes corrupted data on one sector of the disk.



# ALL\_OR\_NOTHING\_PUT

---

1. **procedure** ALMOST\_ALL\_OR\_NOTHING\_PUT (*data, all\_or\_nothing\_sector*)
2.     CAREFUL\_PUT(*data, all\_or\_nothing\_sector.S1*)
3.     CAREFUL\_PUT (*data, all\_or\_nothing\_sector.S2*)     //commit point
4.     CAREFUL\_PUT (*data, all\_or\_nothing\_sector.S3*)
5. **procedure** ALL\_OR\_NOTHING\_GET (**reference** *date, all\_or\_nothing\_sector*)
6.     CAREFUL\_GET (*data1, all\_or\_nothing\_sector.S1*)
7.     CAREFUL\_GET (*data2, all\_or\_nothing\_sector.S2*)
8.     CAREFUL\_GET (*data3, all\_or\_nothing\_sector.S3*)
9.     **if** (*data1 = data2*)   *data*  $\leftarrow$  *data1*
10.    **else** *data*  $\leftarrow$  *data3*

# ALL\_OR\_NOTHING\_PUT

---

1. **procedure** ALL\_OR\_NOTHING\_PUT (*data, all\_or\_nothing\_sector*)
2.     CHECK\_AND\_REPAIR (*all\_or\_nothing\_sector*)
3.     ALMOST\_ALL\_OR\_NOTHING\_PUT (*data, all\_or\_nothing\_sector*)
4. **procedure** CHECK\_AND\_REPAIR (*all\_or\_nothing\_sector*)  
                                  // Ensure copies match
5.     CAREFUL\_GET (*data1, all\_or\_nothing\_sector.S1*)
6.     CAREFUL\_GET (*data2, all\_or\_nothing\_sector.S2*)
7.     CAREFUL\_GET (*data3, all\_or\_nothing\_sector.S3*)

# ALL\_OR\_NOTHING\_PUT

---

8. **if** ( $data1 = data2$ ) **and** ( $data2 = data3$ ) **return** // State 1 or 7, no repair
9. **if** ( $data1 = data2$ )
10.     CAREFUL\_PUT ( $data1, all\_or\_nothing\_sector.S3$ ) **return** // State 5 or 6.
11. **if** ( $data2 = data3$ )
12.     CAREFUL\_PUT ( $data2, all\_or\_nothing\_sector.S1$ ) **return** // State 2 or 3.
13. CAREFUL\_PUT ( $data1, all\_or\_nothing\_sector.S2$ ) // State 4, go to state 5
14. CAREFUL\_PUT ( $data1, all\_or\_nothing\_sector.S3$ ) // State 5, go to state 7

data state:	1	2	3	4	5	6	7
sector $S1$	old	bad	new	new	new	new	new
sector $S2$	old	old	old	bad	new	new	new
sector $S3$	old	old	old	old	old	bad	new

# Atomicity

---

\_\_\_\_\_  
\_\_\_\_\_  
\_\_\_\_\_

**begin all-or-nothing action**

\_\_\_\_\_  
\_\_\_\_\_  
\_\_\_\_\_  
\_\_\_\_\_  
\_\_\_\_\_



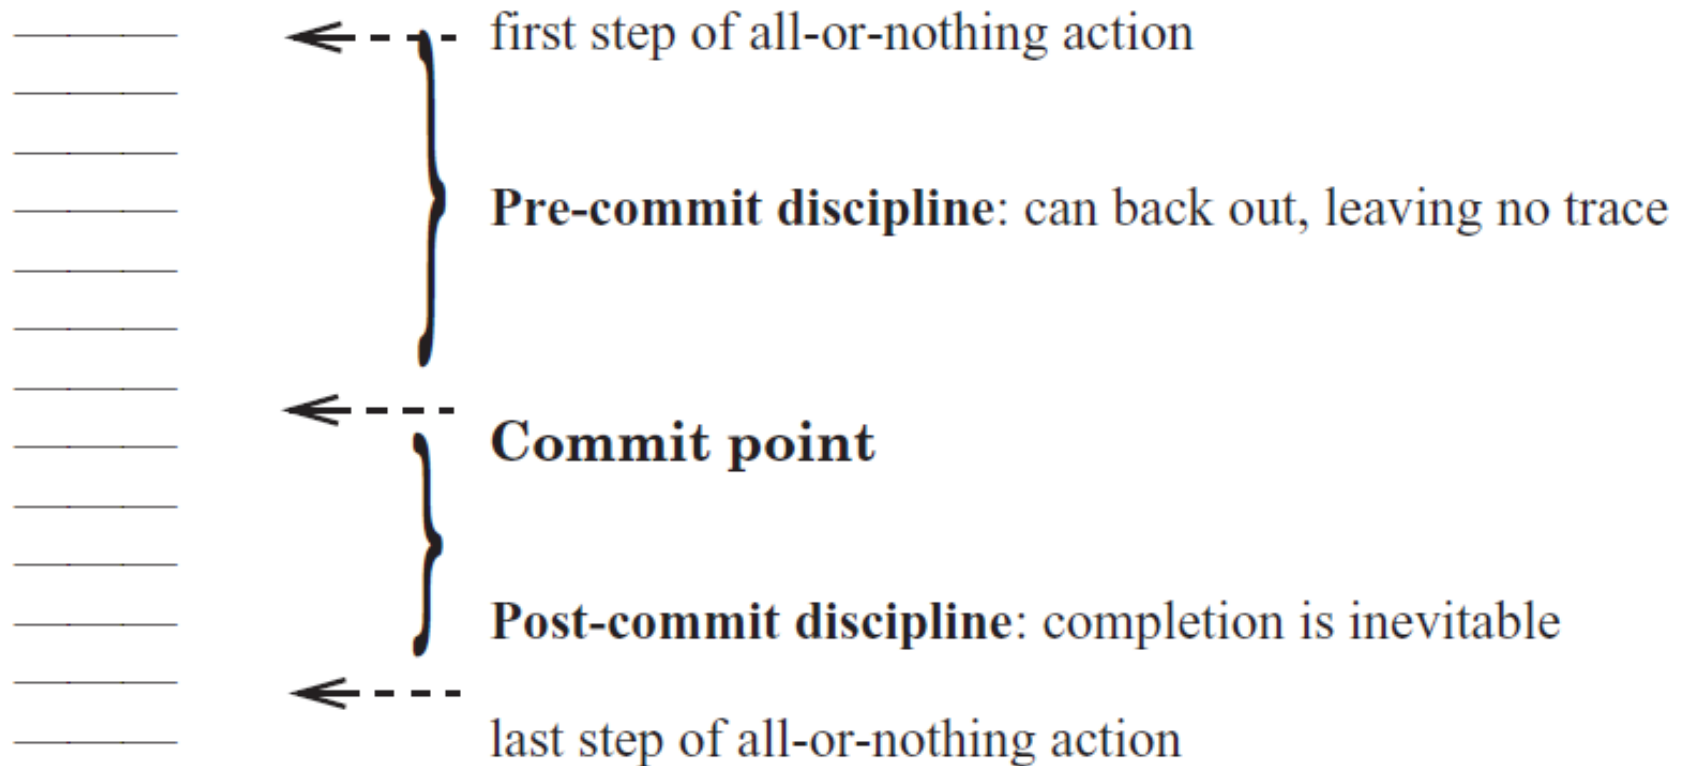
arbitrary sequence of  
lower-layer actions

**end all-or-nothing action**

\_\_\_\_\_  
\_\_\_\_\_  
\_\_\_\_\_

# Commit

---



# Commit

---

- Pre-commit
  - identify all the resources needed to complete the all-or-nothing action,
  - establish their availability
  - maintain the ability to abort at any instant
    - shared resources, once reserved, cannot be released until the commit point is passed
    - should not do anything externally visible
- Post-commit
  - release reserved resources that are no longer needed
  - perform externally visible actions
  - CANNOT try to acquire additional resources
- Q: where's commit in ALL\_OR\_NOTHING\_PUT?

# Shadow Copy

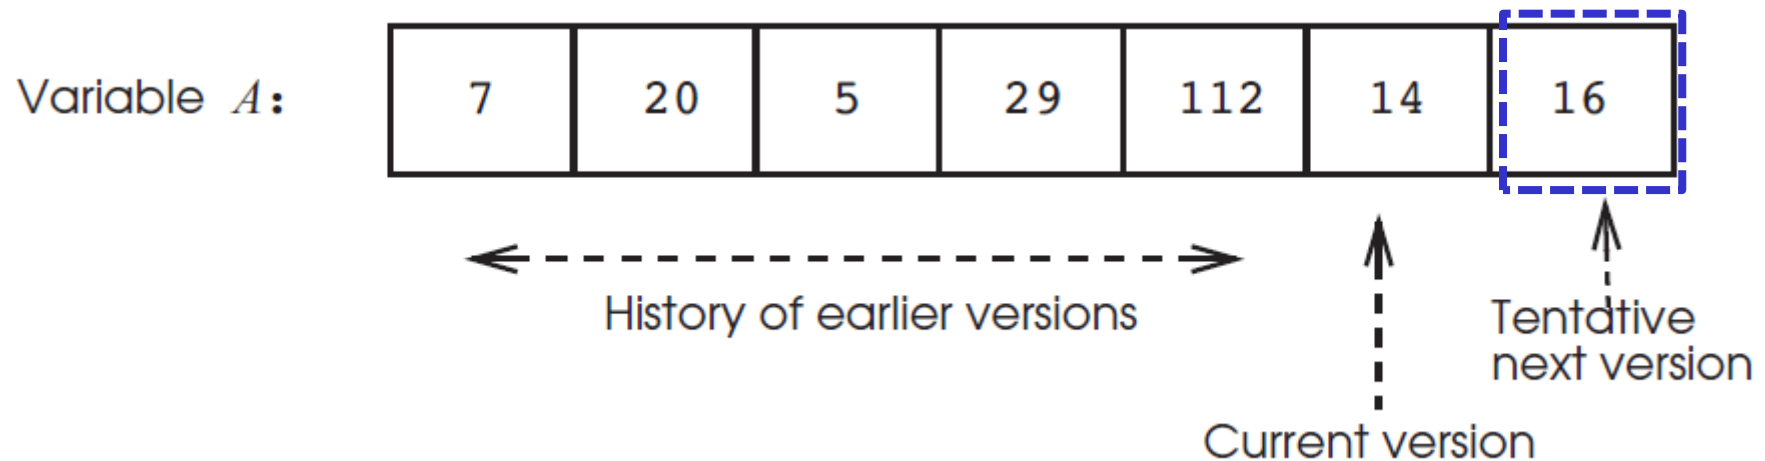
---

- Pre-commit:
  - Create a complete duplicate working copy of the file that is to be modified
  - make all changes to the working copy
- Commit point:
  - Carefully exchange the working copy with the original
  - Typically this step is bootstrapped
- Post-commit:
  - Release the space that was occupied by the original
- The golden rule of atomicity
  - *Never modify the only copy!*

# Journal Storage

---

- a store operation
  - Not overwrites old data
  - create a new, tentative version of the data
    - remains invisible to any reader outside this all-or-nothing action






# Journal Storage

---

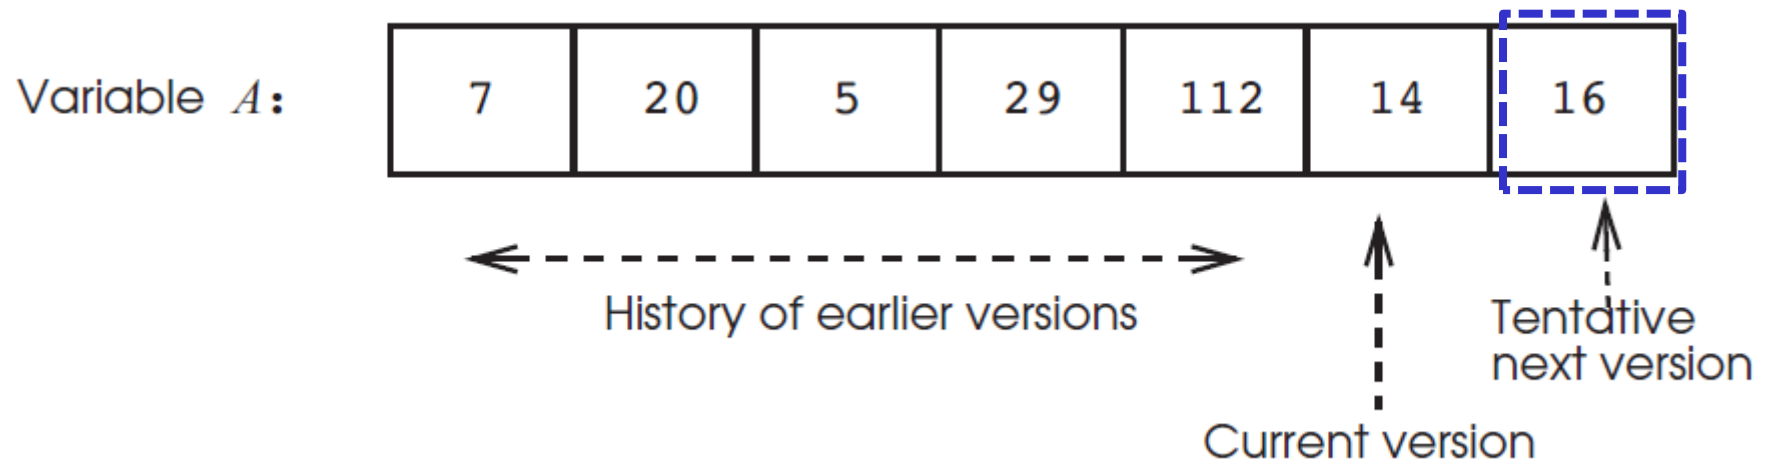
- a store operation
  - Not overwrites old data
  - create a new, tentative version of the data
    - remains invisible to any reader outside this all-or-nothing action



# Journal Storage

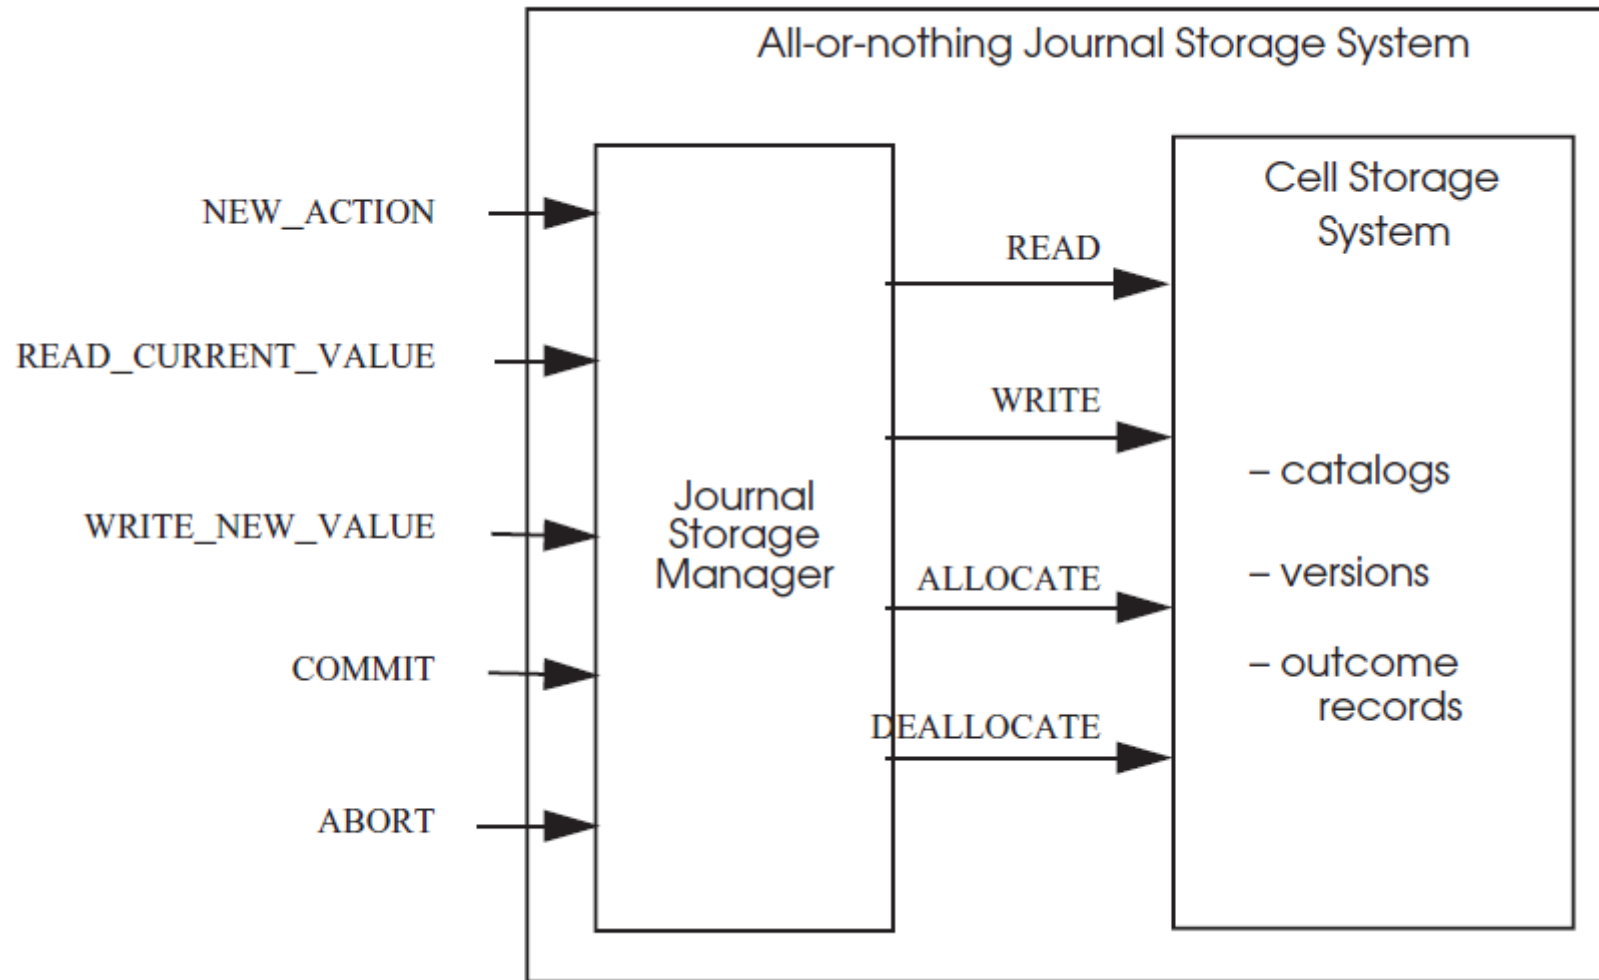
---

- a store operation
  - Not overwrites old data
  - create a new, tentative version of the data
    - remains invisible to any reader outside this all-or-nothing action



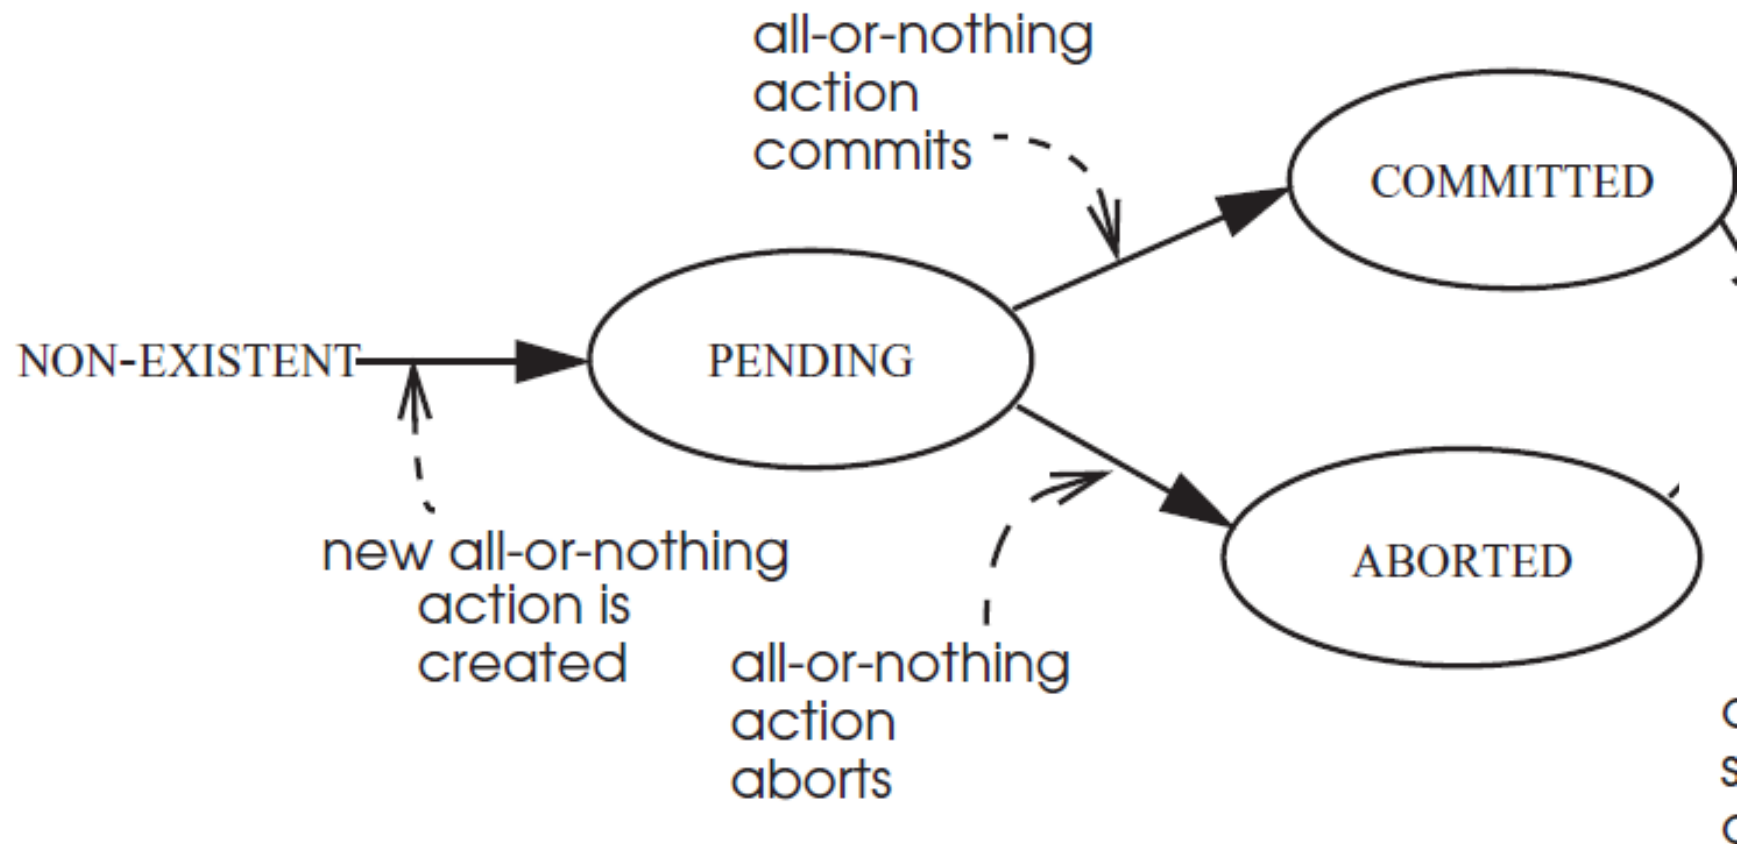
# Journal Storage

---



# Journal Storage

---



# Journal Storage

---

1. **procedure** NEW\_ACTION ()
2.      $id \leftarrow \text{NEW\_OUTCOME\_RECORD } ()$
3.      $id.outcome\_record.state \leftarrow \text{PENDING}$
4.     **return**  $id$
  
5. **procedure** COMMIT (**reference**  $id$ )
6.      $id.outcome\_record.state \leftarrow \text{COMMITTED}$
  
7. **procedure** ABORT (**reference**  $id$ )
8.      $id.outcome\_record.state \leftarrow \text{ABORTED}$

# Journal Storage

---

1. **procedure** READ\_CURRENT\_VALUE (*data\_id*, *caller\_id*)
2.     **starting at end of** *data\_id* **repeat until beginning**
3.         *v* ← **previous version of** *data\_id*   // Get next older version
4.         *a* ← *v.action\_id*                   // Identify the action *a* that created it
5.         *s* ← *a.outcome\_record.state*   // Check action *a*'s outcome record
6.         **if** *s* = **COMMITTED** **then**
7.             **return** *v.value*
8.         **else skip** *v*                                 // Continue backward search
9.     **signal** ("Tried to read an uninitialized variable!")

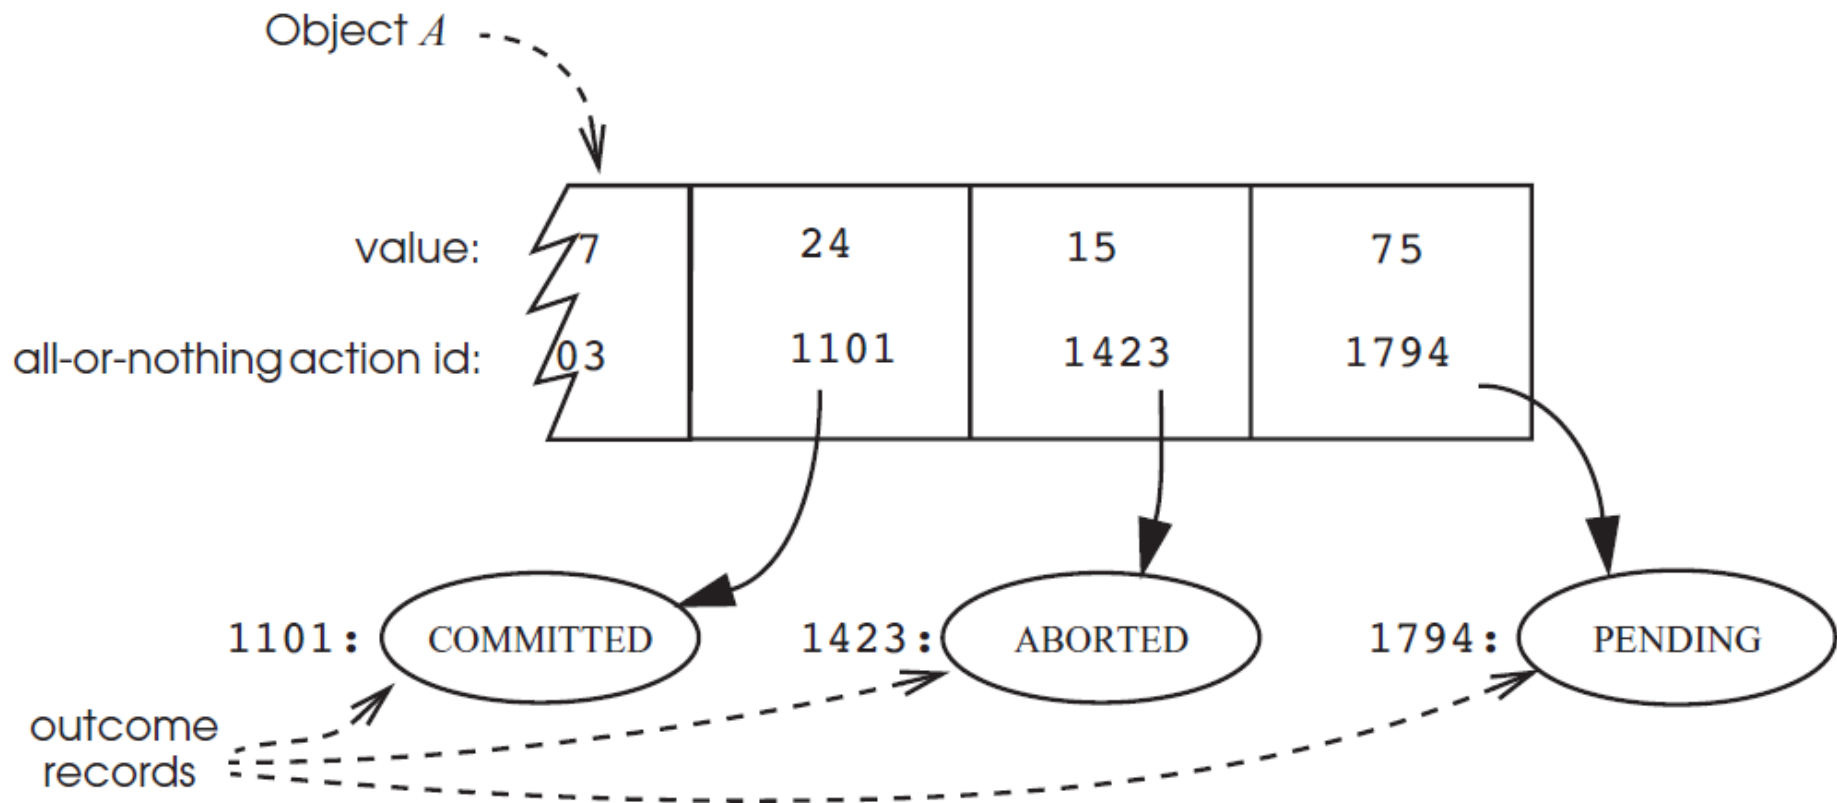
# Journal Storage

---

10. **procedure** WRITE\_NEW\_VALUE (**reference** *data\_id*,  
    *new\_value*, *caller\_id*)
11.     **if** *caller\_id.outcome\_record.state* = **PENDING**
12.         **append new version** *v* **to** *data\_id*
13.         *v.value*  $\leftarrow$  *new\_value*
14.         *v.action\_id*  $\leftarrow$  *caller\_id*
15.     **else signal** (“Tried to write outside of an all-or-nothing action!”)

# Journal Storage

---





# Journal Storage

---

1. **procedure** TRANSFER (**reference** *debit\_account*, **reference** *credit\_account*, *amount*)
2.     *my\_id*  $\leftarrow$  **NEW\_ACTION** ()
3.     *xvalue*  $\leftarrow$  **READ\_CURRENT\_VALUE** (*debit\_account*, *my\_id*)
4.     *xvalue*  $\leftarrow$  *xvalue* - *amount*
5.     **WRITE\_NEW\_VALUE** (*debit\_account*, *xvalue*, *my\_id*)
6.     *yvalue*  $\leftarrow$  **READ\_CURRENT\_VALUE** (*credit\_account*, *my\_id*)
7.     *yvalue*  $\leftarrow$  *yvalue* + *amount*
8.     **WRITE\_NEW\_VALUE** (*credit\_account*, *yvalue*, *my\_id*)
9.     **if** *xvalue* > 0 **then**
10.         **COMMIT** (*my\_id*)
11.     **else**
12.         **ABORT** (*my\_id*)
13.         **signal** (“Negative transfers are not allowed.”)


# Log

---

- The application
  - first logs the change in journal storage
  - then it installs the change in cell storage
- separates the reading and writing of data from the failure recovery mechanism
- Minimize the number of storage accesses required for the most common activities
  - Reads, updates
- Rarely-performed activities may not be minimal
  - failure recovery

# Log

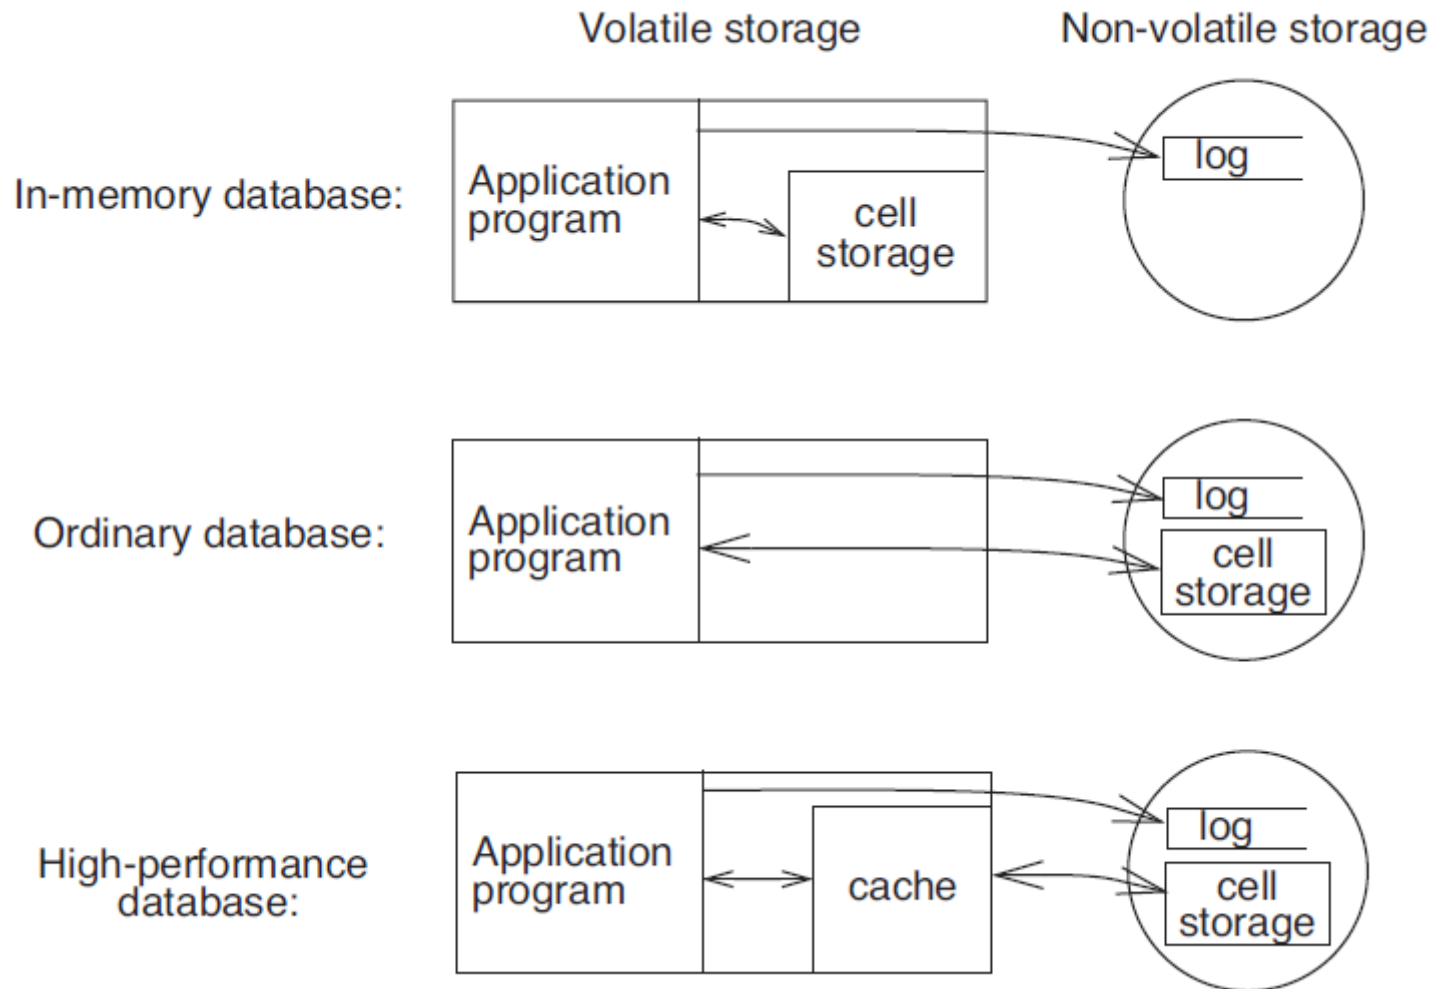
---



**Write-ahead-log protocol (WAL)**  
Log the update before installing it

# Log

---



# Logging Protocol

---

- **CHANGE** Record
  - The identity of the all-or-nothing action
  - (1) A component action redo
    - installs the intended value in cell storage
      - After commit, if the system crashes, the recovery procedure can perform the install on behalf of the action
  - (2) A second component action undo
    - reverses the effect on cell storage of the install
      - After aborts or the system crashes, it may be necessary for the recovery procedure to reverse the effect

# Logging Protocol

---

- The application
- **NEW\_ACTION**
  - log a **BEGIN** record that contains just the new identity
  - As the all-or-nothing action proceeds through its pre-commit phase, it logs **CHANGE** records
- To implement COMMIT or ABORT
  - logs an **OUTCOME** record
  - commit point

# Logging Protocol

---

1. **procedure** **TRANSFER** (*debit\_account, credit\_account, amount*)
2.      $\text{my\_id} \leftarrow \text{LOG}(\text{BEGIN\_TRANSACTION})$
3.      $\text{dbvalue.old} \leftarrow \text{GET}(\text{debit\_account})$
4.      $\text{dbvalue.new} \leftarrow \text{dbvalue.old} - \text{amount}$
5.      $\text{crvalue.old} \leftarrow \text{GET}(\text{credit\_account}, \text{my\_id})$
6.      $\text{crvalue.new} \leftarrow \text{crvalue.old} + \text{amount}$
7.      $\text{LOG}(\text{CHANGE}, \text{my\_id},$
8.          $\text{"PUT (debit\_account, dbvalue.new)"}$ , //redo action
9.          $\text{"PUT (debit\_account, dbvalue.old)"}$  ) //undo action

# Logging Protocol ---transfer (cont.)

---

```
10.  LOG (  CHANGE, my_id,
11.          “PUT (credit_account, crvalue.new)”    //redo action
12.          “PUT (credit_account, crvalue.old)”    ) //undo action
13.  PUT (debit_account, dbvalue.new) // install
14.  PUT (credit_account, crvalue.new) // install
15.  if dbvalue.new > 0 then
16.      LOG ( OUTCOME, COMMIT, my_id)
17.  else
18.      LOG (OUTCOME, ABORT, my_id)
19.      signal(“Action not allowed. Would make debit account negative.”)
20.  LOG (END_TRANSACTION, my_id)
```



# Logging Protocol

---

...	<i>type</i> : CHANGE <i>action_id</i> : 9979 <i>redo_action</i> : PUT( <i>debit_account</i> , \$90) <i>undo_action</i> : PUT( <i>debit_account</i> , \$120)	<i>type</i> : OUTCOME <i>action_id</i> : 9974 <i>status</i> : COMMITTED	<i>type</i> : CHANGE <i>action_id</i> : 9979 <i>redo_action</i> : PUT( <i>credit_account</i> , \$40) <i>undo_action</i> : PUT( <i>credit_account</i> , \$10)
-----	--	---	---

← older log records                      newer log records →

# Logging Protocol

---

1. **procedure** ABORT (*action\_id*)
2.     **starting at end of log repeat until beginning**
3.         *log\_record*  $\leftarrow$  **previous record of log**
4.         **if** *log\_record.id* = *action\_id* **then**
5.             **if** (*log\_record.type* = **OUTCOME**)
6.                 **then signal** (“Can’t abort an already completed action.”)
7.             **if** (*log\_record.type* = **CHANGE**)
8.                 **then perform undo\_action of log\_record**
9.             **if** (*log\_record.type* = **BEGIN**)
10.                 **then break repeat**
11.     LOG (*action\_id*, **OUTCOME**, **ABORTED**) // Block future undos.
12.     LOG (*action\_id*, **END**)

# Logging Protocol: **in-memory** **database**

---

1. **procedure** **RECOVER** () // Recovery procedure for a volatile, in-memory database.
2.     *winners*  $\leftarrow$  NULL
3.     **starting at end of log repeat until beginning**
4.         log\_record  $\leftarrow$  **previous record of log**
5.         **if** (*log\_record.type* = OUTCOME)
6.             **then** *winners*  $\leftarrow$  *winners* + log\_record // Set addition.
7.     **starting at beginning of log repeat until end**
8.         log\_record  $\leftarrow$  **next record of log**
9.         **if** (*log\_record.type* = CHANGE)
10.             **and** (*outcome\_record*  $\leftarrow$  **find** (*log\_record.action\_id*) **in** *winners*)
11.             **and** (*outcome\_record.status* = COMMITTED) **then**
12.             **perform** log\_record.redo\_action

- 
- Soft state
    - Can be discarded, no need to redo
  - Crash during recovery (volatile cell )

# Non-volatile logging

---

- NEW\_ACTION 111

- CHANGE 111

New A → old A

- CHANGE 111

New B → old B

( some install )

- OUTCOME 111

COMMIT

( some install )

- END 111

# Logging Protocol: non-volatile cell memory

```
1  procedure RECOVER ()// Recovery procedure for non-volatile cell memory
2      completeds ← NULL
3      losers ← NULL
4      starting at end of log repeat until beginning
5          log_record ← previous record of log
6          if (log_record.type = END)
7              then completeds ← completeds + log_record           // Set addition.
8              if (log_record.action_id is not in completeds) then
9                  losers ← losers + log_record           // Add if not already in set.
10                 if (log_record.type = CHANGE) then
11                     perform log_record.undo_action

12     starting at beginning of log repeat until end
13         log_record ← next record of log
14         if (log_record.type = CHANGE)
15             and (log_record.action_id.status = COMMITTED) then
16                 perform log_record.redo_action

17     for each log_record in losers do
18         log (log_record.action_id, END)           // Show action completed.
```

Necessary ?

# Undo logging

---

- NEW\_ACTION 114

- CHANGE 114

New A → old A

New B -> old B

( install the cell )

- OUTCOME 114

COMMIT

- END 114

# Undo logging

Logging Protocol: non-volatile cell memory

---

```
1  procedure RECOVER ()           // Recovery procedure for rollback recovery.
2      completeds ← NULL
3      losers ← NULL
4      starting at end of log repeat until beginning           // Perform undo scan.
5          log_record ← previous record of log
6          if (log_record.type = OUTCOME)
7              then completeds ← completeds + log_record           // Set addition.
8          if (log_record.action_id is not in completeds) then
9              losers ← losers + log_record           // New loser.
10             if (log_record.type = CHANGE) then
11                 perform log_record.undo_action

12  for each log_record in losers do
13      log (log_record.action_id, OUTCOME, ABORT)           // Block future undos.
```



## redo logging

---

- NEW\_ACTION 114

- CHANGE 114

New A → old A

New B → old B

- OUTCOME 114

COMMIT

( install the cell )

- END 114

#1067, #1081, #1082: checkpoint

#1083: start

#1084: start

#1083: set y 5 -> 6

#1083: set x 5 -> 9

#1083: commit

#1084: set y 6 -> 4

#1085: start

#1085: set z 3 -> 4

#1067: abort

#1081: set q 1 -> 9

#1086: start

#1085: set y 4 -> 3

#1084: commit

#1085: set y 3 -> 7

#1081: commit

#1087: start

#1086: set x 9 -> 2

#1086: set w 0 -> 1

#1086: commit

#1087: set u 2 -> 1

===== ⬅ FATAL SYSTEM ERROR!

#1067, #1081, #1082: checkpoint

#1083: start

#1084: start

#1083: set y 5 -> 6

#1083: set x 5 -> 9

#1083: commit

#1084: set y 6 -> 4

#1085: start

#1085: set z 3 -> 4

#1067: abort

#1081: set q 1 -> 9

#1086: start

#1085: set y 4 -> 3

#1084: commit

#1085: set y 3 -> 7

#1081: commit

#1087: start

#1086: set x 9 -> 2

#1086: set w 0 -> 1

#1086: commit

#1087: set u 2 -> 1

===== ⬅ FATAL SYSTEM ERROR!

---

**winners: 1067, 1081, 1084,  
1083, and 1086**

**losers: 1082, 1085, and 1087**