



High-dimensional & Sparse Data

- In real-world intelligent systems, data are usually high-dimensional and very sparse
 - ▣ Text data: Bag-of-words representation
 - ▣ Image data: Visual dictionary representation
 - ▣ Rating data: Sparse adjacent matrix representation
 - ▣ etc.

Google Baidu 百度 Bing

An Example of Text Data Representation

- Each text (e.g., document, paper, webpage, etc.) can be represented by a set of n -grams
- For example, a sentence “This is a short sentence”
 - $n = 1$: { “This” , “is” , “a” , “short” , “sentence” }
 - $n = 2$: { “This is” , “is a” , “a short” , “short sentence” }
 - $n = 3$: { “This is a” , “is a short” , “a short sentence” }
 - In practice, it is common to adopt $n \geq 5$
- Using n -grams will lead to extremely high-dimensional feature vectors: $D = (10^5)^5 = 10^{25} = 2^{83}$
- In current practice, $D = 2^{64}$ seems sufficient

Computation of Similarity

- Computation of data distance is essential in machine learning and thus big data analytics
 - Nearest Neighbor (NN) search for retrieval

$$x^* = \arg \min_{x_n \in S} ||x_q - x_n||^2$$
 - Similarity (Distance) based clustering

$$\min \frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K z_{n,k} ||x_n - \mu_k||^2$$
- Computation of inner product is common in linear kernel methods
 - Support Vector Machine
 - Gaussian Process
 - Kernel PCA

Linear Kernel Methods

- Recall the objective function for a linear regression

$$\min_w \frac{1}{N} \sum_{n=1}^N (y_n - w^\top x_n)^2 + \lambda \|w\|^2$$

- Set the derivatives of the objective w.r.t. w to zero

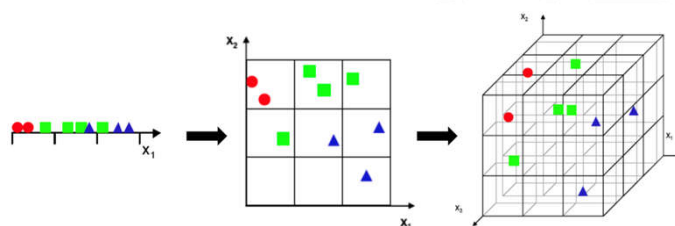
$$\begin{aligned} \frac{\partial J(w)}{\partial w} &= \frac{2}{N} \sum_{n=1}^N (w^\top x_n - y_n) x_n + 2\lambda w = 0 \\ \Rightarrow w &= \frac{1}{\lambda N} \sum_{n=1}^N (y_n - w^\top x_n) x_n = \sum_{n=1}^N \alpha_n x_n \end{aligned}$$

- Substitute $w = \sum_{n=1}^N \alpha_n x_n$ into the objective function we can obtain the dual representation

$$\min_{\alpha} (Y - \sum_{n=1}^N \alpha_n x_n)^\top (Y - \sum_{n=1}^N \alpha_n x_n) + \lambda \sum_{n=1}^N \alpha_n^2$$

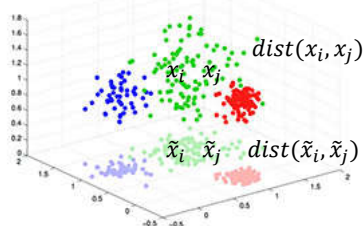
Challenges with High-dimensional Data

- High computational cost: Increase matrix operations and increase search space
- High storage cost: Too large to store high-dimensional data and difficult to load the big model into memory
- The curse of dimensionality:** The volume of the space increases so fast that the available data become sparse



Random Projection

- Recall that dimensionality reduction methods (e.g., PCA) are themselves learning based
- Are there methods that is able to generate a projection matrix without learning and satisfies $\text{dist}(x_i, x_j) \approx \text{dist}(\tilde{x}_i, \tilde{x}_j)$
 - $\tilde{x}_n = H^T x_n$, where $H \in R^{D \times d}$ ($d \ll D$) is a projection matrix
 - $\text{dist}(\cdot, \cdot)$ is a distance function (or similarity measure)



Random Projection: JL Lemma

- **Johnson-Lindenstrauss Lemma**^[1]: Given $0 < \epsilon < 1$, a set X of N points in R^D , and a number $d > 8 \ln N / \epsilon^2$, there exists a linear mapping $H: R^D \rightarrow R^d$ such that for all $x_i, x_j \in X$

$$(1 - \epsilon) \|x_i - x_j\|^2 \leq \|Hx_i - Hx_j\|^2 \leq (1 + \epsilon) \|x_i - x_j\|^2$$

- The JL-lemma states that a small set of points in a high-dimensional space can be embedded into a space of much lower dimension in such a way that **distances between the points are nearly preserved**.
- The JL-Lemma also holds for dot products

$$x_i^T x_j - \epsilon \leq (Hx_i)^T Hx_j \leq x_i^T x_j + \epsilon$$

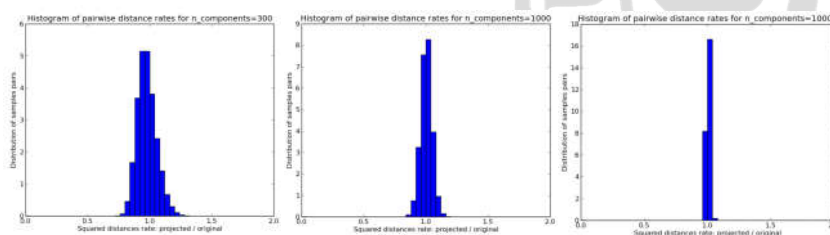
[1] Johnson & Lindenstrauss (1984). "Extensions of Lipschitz mappings into a Hilbert space".

The Random Projection Algorithm

- Random Projection: **Fast, efficient and distance-preserving dimensionality reduction technique!**
- The core idea behind random projection is given in the Johnson-Lindenstrauss lemma
- Canonical (Gaussian) Random Projection:
 - Construct a random matrix $H' \in R^{D \times d}$ ($d \ll D$) by picking the entries from a univariate Gaussian $N(0, \sigma^2)$
 - Orthonormalize the rows of H' to get H
 - Project a data point x_n in the original D -dimensional space into the new d -dimensional space: $\tilde{x}_n = H^T x_n$

Random Projection Example

- Apply Gaussian random projection to the 20-Newsgroups dataset with different configuration of dimensions
 - From 100.000 features to 300 (0.3%)
 - From 100.000 features to 1.000 (1%)
 - From 100.000 features to 10.000 (10%)



[1] <https://www.cs.toronto.edu/~duvenaud/talks/random-kitchen-sinks-tutorial.pdf>

Document Classification

- In a document classification task, the input to the machine learning algorithm is raw text, represented by a bag of words (BoW) representation:
 - ❑ The individual tokens are extracted and counted, and each distinct token in the training set defines a feature.
 - ❑ The BoW for a set of documents is regarded as a term-document matrix where each row is a single document, and each column is a single feature (word).
 - ❑ The entry i, j in such a matrix captures the frequency (or weight) of the j th term of the vocabulary in document i .
 - ❑ The common approach is to construct a dictionary of the training set, and use that to map words to indices.

Document Classification

- An example of BoW representation of documents
 - ❑ Document 1: "He studies machine learning".
 - ❑ Document 2: "Machine learning is interesting".
 - ❑ Document 3: "Machine learning supports big data".

	Dictionary								
	He	studies	machine	learning	is	interesting	supports	big	data
Doc1	1	1	1	1	0	0	0	0	0
Doc2	0	0	1	1	1	1	0	0	0
Doc3	0	0	1	1	0	0	1	1	1

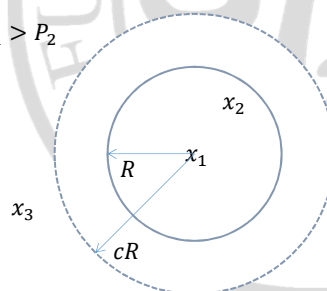
- The problem with this process is that such dictionaries take up a large amount of storage space and **grow in size as the training set grows**.

Locality-Sensitive Hashing

- Locality-Sensitive Hashing (LSH) hashes input items so that similar items map to the same "buckets" with high probability
 - If $d(x_1, x_2) \leq R$, then $h(x_1) = h(x_2)$ with high probability at least P_1
 - If $d(x_2, x_3) \geq cR$, then $h(x_2) = h(x_3)$ with low probability at most P_2
 - An LSH family is interesting only $P_1 > P_2$

- Alternative definition of LSH

$$E_h[h(x_i) = h(x_j)] = \text{sim}(x_i, x_j)$$



LSH: SimHash

- SimHash is designed to approximate the **cosine similarity** $\cos(\theta(x_i, x_j))$ between vectors x_i and x_j .
- SimHash is used by the Google Crawler to find near duplicate pages
- Given an input vector x_i and a random hyperplane specified by a normal unit vector r , the SimHash function is defined as $h(x_i) = \text{sgn}(r^\top x_i)$
- Randomly choose multiple hyperplanes and the limit of the collision ratio equals to the probability of hyperplane falling in the angle between the two vectors $\frac{\theta(x_i, x_j)}{\pi}$

$$\frac{1}{N} \sum_k 1(h_k(x_i) = h_k(x_j)) \xrightarrow{k \rightarrow \infty} E_h[h(x_i) = h(x_j)] = 1 - \frac{\theta(x_i, x_j)}{\pi}$$

[1] Charikar (2002), Similarity Estimation Techniques from Rounding Algorithms.

LSH: MinHash

- MinHash is designed to approximate the **Jaccard similarity**

$$J(S_i, S_j) = \frac{|S_i \cap S_j|}{|S_i \cup S_j|} \text{ between sets } S_i \text{ and } S_j.$$

- MinHash is the first position of element after a random permutation $h(S_i) = \min(\pi(S_i))$.
- Property of MinHash

$$J(S_i, S_j) = \frac{|S_i \cap S_j|}{|S_i \cup S_j|} = E_{\pi}[1(\min(\pi(S_i)) = \min(\pi(S_j)))]$$

- Empirically use K independent random permutations for approximation

$$J(S_i, S_j) \approx \frac{1}{K} \sum_{k=1}^K 1(\min(\pi_k(S_i)) = \min(\pi_k(S_j)))$$

MinHash Example

- MinHash based text classification
 - Represent each text as a bag-of-words
 - Compute pairwise Jaccard similarities based on MinHashes
 - Construct the kernel matrix



$$h_1(S_1) = 1 \quad h_2(S_1) = 1 \quad h_3(S_1) = 2 \quad h_4(S_1) = 2$$

Machine Learning Journal Conference Research	Conference Learning Machine Research Journal	Journal Machine Learning Research Conference	Research Conference Machine Learning Journal
--	--	--	--

$$h_1(S_2) = 1 \quad h_2(S_2) = 2 \quad h_3(S_2) = 1 \quad h_4(S_2) = 1$$

$$J(S_1, S_2) = \frac{1}{4}$$

Summary

- Random Projection $\rightarrow E[\langle Hx_i, Hx_j \rangle] = \langle x_i, x_j \rangle$
 - Johnson-Lindenstrauss transform
 - Gaussian random projection
 - etc.
- Locality-Sensitive Hashing $\rightarrow E_h[h(x_i) = h(x_j)] = \text{sim}(x_i, x_j)$
 - SimHash
 - MinHash
 - etc.

Kernel Methods

- We can efficiently embed high-dimensional data satisfying $E[\langle Hx_i, Hx_j \rangle] = \langle x_i, x_j \rangle$ or $E_h[h(x_i) = h(x_j)] = \text{sim}(x_i, x_j)$ now.
- Recall the linear kernel methods introduced before
 - The classifier $f(x) = w^\top x = \sum_{n=1}^N \alpha_n x_n^\top x$
 - The objective $\min_{\alpha} (Y - XX^\top \alpha)^\top (Y - XX^\top \alpha) + \lambda \alpha^\top XX^\top \alpha$
- Only the **inner product** appears in both the objective function and the classifier
 - No need to know the exact form of x
 - The inner product can be replaced by any (approximate) similarity measure

Kernel Methods

- Replace the inner product by the kernel function
 - The classifier becomes $f(x) = w^\top x = \sum_{n=1}^N \alpha_n \kappa(x_n, x)$
 - The objective becomes $\min_{\alpha} (Y - K\alpha)^\top (Y - K\alpha) + \lambda \alpha^\top K \alpha$, where $K_{i,j} = \kappa(x_i, x_j)$
- Now we only need to let the kernel function be
 - Random projection: $\kappa(x_i, x_j) = \langle Hx_i, Hx_j \rangle$
 - SimHash: $\kappa(x_i, x_j) = \frac{1}{K} \sum_{k=1}^K 1(h_k(x_i) = h_k(x_j))$
 - MinHash: $\kappa(x_i, x_j) = \frac{1}{K} \sum_{k=1}^K 1(\min(\pi_k(S_i)) = \min(\pi_k(S_j)))$
 - etc.
- A valid kernel – The kernel matrix K must be positive semi-definite



Thanks

Email: libin@fudan.edu.cn