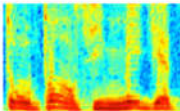




## Why Word Embedding

- Natural language processing systems traditionally treat words as discrete atomic symbols, provide no useful information to the system regarding the relationships that may exist between the individual symbols

AUDIO



Audio Spectrogram

DENSE

IMAGES


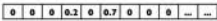


Image pixels

DENSE

TEXT



Word, context, or document vectors

SPARSE

## One-Hot Representation

- Represent a word as a **one-hot** vector
  - Example: He studies machine learning

Dictionary									
	He	studies	machine	learning	is	interesting	supports	big	data
$v_{He}$	1	0	0	0	0	0	0	0	0
$v_{is}$	0	0	0	0	1	0	0	0	0
$v_{big}$	0	0	0	0	0	0	0	1	0
$v_{data}$	0	0	0	0	0	0	0	0	1

- How large is this dictionary (universe set)?
  - Penn Treebank dataset: ~50K
  - Google 1T dataset: 13M

## Issues of One-Hot Vector

- High-dimensional
- Sparse
- Fixed dimensionality (cannot represent new words)
- Orthogonal semantic similarity between pair of words

Dictionary							
	king	queen	professor	interesting	supports	big	data
$v_{king}$	1	0	0	0	0	0	0
$v_{queen}$	0	1	0	0	0	0	0
$v_{professor}$	0	0	1	0	0	0	0

$$\langle v_{king}, v_{queen} \rangle = \langle v_{king}, v_{professor} \rangle = 0$$

## Distributional Representation

- "You shall know a word by the company it keeps" (John R. Firth, 1957)
- A word is characterized by its context

Dictionary						
	royal	palace	duke	speech	university	research
$v_{king}$	1	1	1	1	0	0
$v_{queen}$	1	1	1	1	0	0
$v_{professor}$	0	0	0	1	1	1

$$\langle v_{king}, v_{queen} \rangle > \langle v_{king}, v_{professor} \rangle = \langle v_{queen}, v_{professor} \rangle$$

- Still not good enough ...

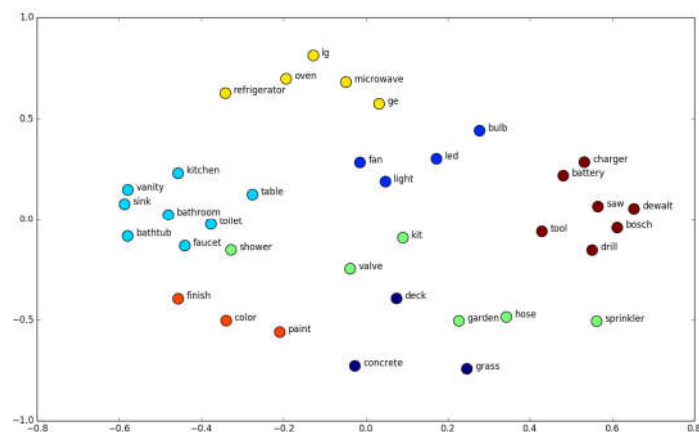
## Vector Representation

- The vector space is spanned by semantic "concepts"
- Each word is represented by a distribution of weights over these concepts
  - The representation of a word is spread across all of the concepts in the vector
  - Each concept in the vector contributes to the definition of many words

Concepts				
	Royalty	Masculinity	Femininity	Celebrity
$v_{king}$	0.9	0.9	0.1	0.9
$v_{queen}$	0.8	0.2	0.9	0.8
$v_{actor}$	0.1	0.8	0.2	0.7

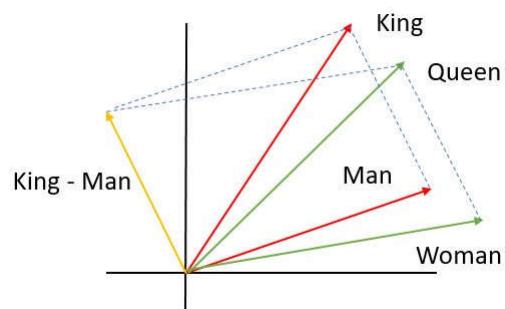
## Vector Representation

- An illustration of 2-D vector representation



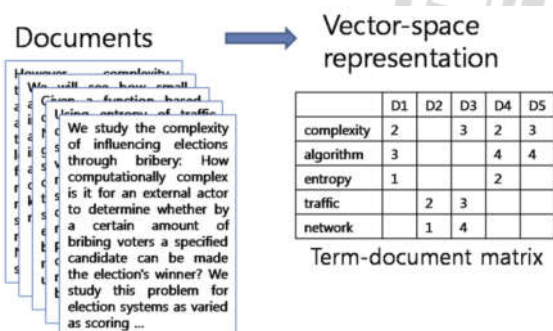
## How to Learn Word Vectors?

- How to find semantic concepts - **bases**
- How to assign weights - **vectors**
- How to define similarity/distance - **metric**



## A Simple Vector Representation

- A word is represented by the documents (bases of the vector space) in which it appears
- A document is represented by the words it contain (i.e., bag-of-words representation for the document)

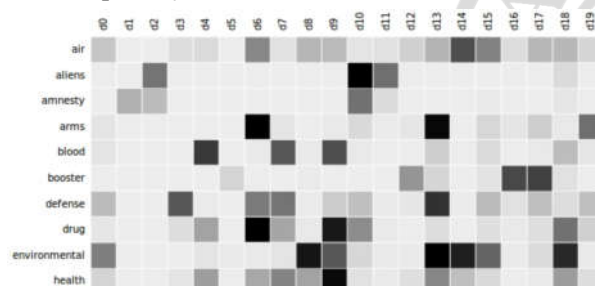


## Issues of Doc-Word Co-occurrence

- Number of concepts (bases) is too large
- Concepts (bases) are not orthogonal
- High-dimensional
- Sparse
- etc.

## Latent Semantic Analysis

- Represent a corpus as a document-word co-occurrence matrix (frequency, tf-idf, etc.)



- Factorize the document-word co-occurrence matrix to find latent components (semantic concepts)

## Latent Semantic Analysis

- Latent semantic analysis (LSA) is a technique of analyzing relationships between a set of documents and the terms they contain by **producing a set of concepts** related to the documents and terms.
- LSA assumes that words that are close in meaning will occur in similar documents (the distributional hypothesis).
- LSA applies singular value decomposition (SVD) to find latent concepts  $A = USV^T$
- Words are then compared by taking the cosine of the angle between the two vectors.

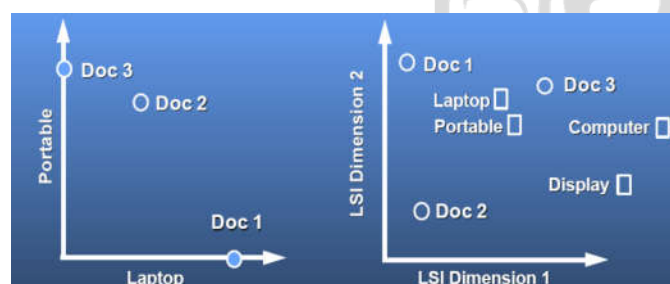
## Latent Semantic Analysis

- LSA applies singular value decomposition (SVD) to find latent concepts  $A = USV^T$ 
  - $A$ :  $m \times n$  word-document co-occurrence matrix
  - $U$ :  $m \times k$  orthogonal matrices for representing words
  - $V$ :  $n \times k$  orthogonal matrices for representing documents
  - $S$ :  $k \times k$  diagonal singular value matrix
  - Select  $k' \ll n, k' \ll m$  for a low-rank approximation of  $A$

A					=	U					x	S					x	Vt				
	d1	d2	d3	d4			f1	f2	f3	f4			f1	f2	f3	f4			d1	d2	d3	d4
a	6	7	1	0		a	0.24	-0.51	0.08	0.06		f1	23.1	0	0	0		f1	0.37	0.38	0.65	0.53
b	8	6	0	1		b	0.25	-0.54	-0.64	-0.23		f2	0	14.3	0	0		f2	-0.55	-0.63	0.37	0.38
c	6	9	8	5		c	0.58	-0.28	0.57	0.13		f3	0	0	3.5	0		f3	-0.69	0.59	0.27	-0.21
d	0	1	8	8		d	0.42	0.37	0.16	-0.68		f4	0	0	0	1.5		f4	0.26	-0.29	0.59	-0.69
e	2	0	9	7		e	0.44	0.34	-0.24	0.66												
f	2	0	7	7		f	0.39	0.29	-0.40	-0.09												

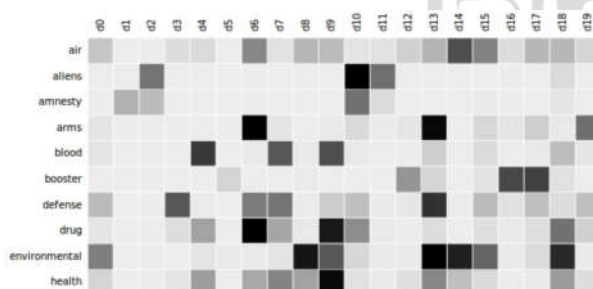
## Latent Semantic Analysis

- After applying SVD to the word-document co-occurrence matrix and obtain the factorization  $A = USV^T$ 
  - $U$ : similar words have large inner products
  - $V$ : similar documents have large inner products
  - Related word and document have large inner products



## Probabilistic LSA

- Probabilistic LSA (PLSA) is a statistical technique for the analysis of co-occurrence matrix.
- Compared to standard LSA stemming from a **low-rank decomposition** (SVD), PLSA is based on a **mixture decomposition** derived from a latent class model



## PLSA Model

- Observations in the form of co-occurrences  $(w, d)$  of words and documents
- PLSA models the probability of  $(w, d)$  as a mixture of conditionally independent multinomial distributions

$$p(w, d) = \sum_z p(z)p(d|z)p(w|z) = p(d) \sum_z p(z|d)p(w|z)$$

- An advantage of PLSA is that the latent variable  $z$  can be interpreted as a topic
  - $z$ : topic (latent class)
  - $p(z|d)$ : each document has a distribution over  $K$  latent topics
  - $p(w|z)$ : each topic has a distribution over the vocabulary

[1] Hofmann (1999). "Probabilistic Latent Semantic Indexing".

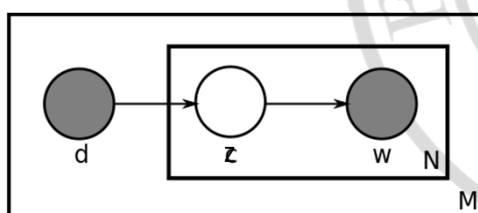


## PLSA Model

- PLSA is a generative model of the documents in the collection it is estimated on

$$p(w, d) = p(d) \sum_z p(z|d)p(w|z)$$

- For each document  $d$ , a topic  $z$  is generated conditionally to  $d$  according to  $p(z|d)$
- A word is then generated from topic  $z$  according to  $P(w|z)$



## EM Algorithm for Latent Variable Models

- Given a joint distribution  $p(X, Z|\theta)$  over observed variables  $X$  and **latent variables**  $Z$ , governed by **parameters**  $\theta$ , the goal is to maximize the likelihood function  $p(X|\theta)$  w.r.t.  $\theta$ .
- The general EM algorithm:
  - Initialize the parameters  $\theta^{\text{old}}$ ;
  - E-Step: Evaluate  $p(Z|X, \theta^{\text{old}})$ ;
  - M-Step: Evaluate  $\theta^{\text{new}}$  given by

$$\theta^{\text{new}} = \operatorname{argmax}_{\theta} \sum_z p(Z|X, \theta^{\text{old}}) p(X, Z|\theta)$$

- Check the convergence of the parameter values; if not convergence condition not satisfied set  $\theta^{\text{old}} = \theta^{\text{new}}$  and go to E-step.

## Learning for PLSA

- The parameters  $p(z|d)$  and  $p(w|z)$  of PLSA can be learned by using the EM algorithm

- EM algorithm for PLSA:

- E-Step: Evaluate  $p(z_k | d_i, w_j; \theta^{\text{old}})$ ;

$$p(z_k | d_i, w_j; \theta^{\text{old}}) = \frac{p(w_j | z_k) p(z_k | d_i)}{\sum_{l=1}^K p(w_j | z_l) p(z_l | d_i)}$$

- M-Step: Evaluate  $\theta^{\text{new}}$  given by

$$p(w_j | z_k) = \frac{\sum_{i=1}^M \#(d_i, w_j) p(z_k | d_i, w_j)}{\sum_{n=1}^N \sum_{m=1}^M \#(d_m, w_n) p(z_k | d_m, w_n)}$$

$$p(z_k | d_i) = \frac{\sum_{j=1}^N \#(d_i, w_j) p(z_k | d_i, w_j)}{\#(d_i)}$$

## word2vec

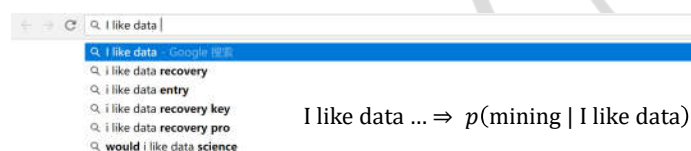
- Latent semantic analysis (LSA)
  - Low-rank factorization of the **co-occurrence** matrix
  - Latent space can be interpreted as **latent concepts**
  - Words are distributional representations in the latent space
- word2vec
  - Word vectors are positioned in the vector space such that words that **share common contexts** in the corpus are located in close proximity to one another in the space
  - Predict surrounding words (skip-gram)
  - Also can be used represent similarity

[1] Mikolov (2013). "Efficient Estimation of Word Representations in Vector Space".

## Language Model

- A statistical language model is a probability distribution over sequences of words  $w_1, \dots, w_N$
- Given such a sequence, it assigns a probability  $p(w_1, \dots, w_N)$  to the whole sequence
  - Rank possible sentences (e.g., spelling correction)
 
$$p(\text{"I like data analytics"}) > p(\text{"I like Dota analytics"})$$

$$p(\text{"I like data analytics"}) > p(\text{"Data analytics like I"})$$
  - Generate possible sentences (e.g., autocomplete query)



## $n$ -gram Language Model

- The probability of a word only depends on the previous  $n - 1$  words, known as an  $n$ -gram model

$$p(w_1, \dots, w_N) = \prod_{i=1}^N p(w_i | w_1, \dots, w_{i-1}) \approx \prod_{n=1}^N p(w_i | w_{i-(n-1)}, \dots, w_{i-1})$$

- Bigram ( $n = 2$ ) language model

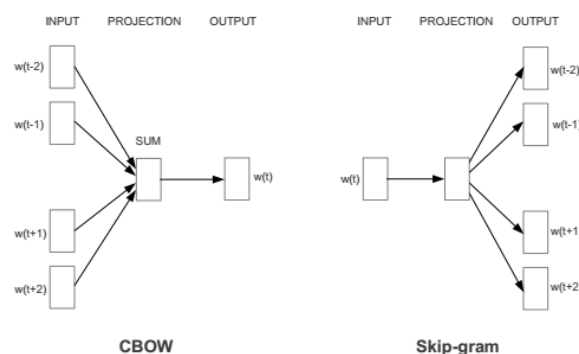
$$p(\text{"I like data analytics"}) \approx p(I | \langle s \rangle) p(\text{like} | I) p(\text{data} | \text{like}) p(\text{analytics} | \text{data}) p(\langle /s \rangle | \text{analytics})$$

- The conditional probability can be calculated from  $n$ -gram model frequency counts

$$p(w_i | w_{i-(n-1)}, \dots, w_{i-1}) = \frac{\#(w_{i-(n-1)}, \dots, w_{i-1}, w_i)}{\#(w_{i-(n-1)}, \dots, w_{i-1})}$$

## CBOW and Skip-Grams

- word2vec can use either continuous bag-of-words (CBOW) or continuous skip-gram to produce a distributed representation of words



## Objective of word2vec (Skip-gram)

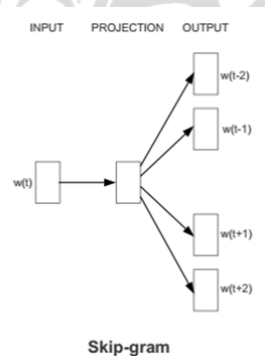
- Maximize the log likelihood of the context words  $w_{t-m}, w_{t-m+1}, \dots, w_{t-1}, w_{t+1}, w_{t+2}, \dots, w_{t+m}$ , given  $w_t$

□  $m$  is usually 5~10

$$J(\theta) = \frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log p(w_{t+j} | w_t)$$

- Use softmax to model  $p(w_{t+j} | w_t)$

$$p(w_{t+j} | w_t) = \frac{\exp(v_{w_{t+j}} \cdot v_{w_t})}{\sum_{w' \in \text{Context}} \exp(v_{w'} \cdot v_{w_t})}$$



## Optimization of word2vec

- How to minimize the objective of word2vec to obtain  $v_{w_t}$  for  $w_1, \dots, w_T$ ? - Gradient descent

- Let the current center word be  $c$  and one of its context word be  $s$ , then the conditional probability becomes

$$p(s|c) = \frac{\exp(v_s \cdot v_c)}{\sum_{w'} \exp(v_{w'} \cdot v_c)}$$

- The gradient of the log likelihood w.r.t.  $v_c$  is

$$\frac{\partial \log p(s|c)}{\partial v_c} = v_s - \sum_{w'} \frac{\exp(v_{w'} \cdot v_c)}{\sum_{w''} \exp(v_{w''} \cdot v_c)} v_{w'} = v_s - E_{w' \sim p(w'|c)} v_{w'}$$

- Alternate minimize  $J(\theta)$  w.r.t.  $v_{w_t}$  for  $w_1, \dots, w_T$

## Optimization of word2vec

- Gradient descent

- Let  $J(\theta) = \frac{1}{n} \sum_{i=1}^n J_i(\theta)$
  - update rule:  $\theta \leftarrow \theta - \frac{\eta}{n} \sum_{i=1}^n \nabla J_i(\theta)$

- Stochastic gradient descent

- Replace  $\frac{1}{n} \sum_{i=1}^n \nabla J_i(\theta)$  by the gradient at a single example  $\nabla J_i(\theta)$
  - At each iteration **randomly select** an example  $i$  and update:  
 $\theta \leftarrow \theta - \eta \nabla J_i(\theta)$

