

Simple File System

i. Project Goals

The goals of this project are:

- to implement a simple file system on top of a virtual disk
- to understand implementation details of file systems

ii. Implementing a simple file system

The goal of this project is to implement a simple file system on top of a virtual disk. To this end, you will implement a library that offers a set of basic file system calls (such as open, read, write, ...) to applications. The file data and file system meta-information will be stored on a virtual disk. This virtual disk is actually a single file that is stored on the "real" file system provided by the Linux operating system. That is, you are basically implementing your file system on top of the Linux file system.

To create and access the virtual disk, we have provided a few definitions and helper functions that you can find in this *header* file and this *source* file. **Note that**, in your library, you are not allowed to create any "real" files on the Linux file system itself. Instead, you have to use the provided helper functions and store all the data that you need on the virtual disk. As you can see by looking at the provided header and source files, the virtual disk has 8,192 blocks, and each block holds 4KB. You can create an empty disk, open and close a disk, and read and write entire blocks (by providing a block number in the range between 0 and 8,191 inclusive).

To make things easier, your file system does not have to support a directory hierarchy. Instead, all files are stored in a single root directory on the virtual disk. In addition, your file system does not have to store more than 64 files (of course, you can create and delete files, and deleted files do not count against this 64 file limit). Finally, out of the 8,192 blocks available on disk, only 4,096 must be reserved as data blocks. That is, you have ample of space to store your meta-information. However, you have to free data blocks (make them available again) when the corresponding file is deleted. The maximum file size is 16 megabytes (all 4,096 data blocks, each with 4KB).

To manage your file system, you have to provide the following three functions:

File system management		
<pre>int make_fs(char *disk_name);</pre>	<p>Create a fresh (and empty) file system on the virtual disk with name <i>disk_name</i>.</p> <p>Return 0 on success, and -1 when the disk <i>disk_name</i> could not be created, opened, or properly initialized.</p> <p>As part of this function, you should first invoke <i>make_disk(disk_name)</i> to create a new disk. Then, open this disk and write/initialize the necessary meta-information for your file system so that it can be later used (mounted).</p>	<p>30' (Required, each for 10')</p>
<pre>int mount_fs(char *disk_name);</pre>	<p>Mount a file system that is stored on a virtual disk with name <i>disk_name</i>.</p> <p>Return 0 on success, and -1 when the disk <i>disk_name</i> could not be opened or when the disk does not contain a valid file system (that you previously created with <i>make_fs</i>).</p> <p>With the mount operation, a file system becomes "ready for use." You need to open the disk and then load the meta-information that is necessary to handle file system operations.</p>	

<pre>int umount_fs(char *disk_name);</pre>	<p>Unmount file system from a virtual disk with name <i>disk_name</i>.</p> <p>Return 0 on success, and -1 when the disk <i>disk_name</i> could not be closed or when data could not be written to the disk (this should not happen).</p> <p>As part of this operation, you need to write back all meta-information so that the disk persistently reflects all changes that were made to the file system (such as new files that are created, data that is written, ...). You should also close the disk.</p>	
--	--	--

It is important to observe that your file system must provide **persistent storage**. That is, assume that you have created a file system on a virtual disk and mounted it. Then, you create a few files and write some data to them. Finally, you unmount the file system. At this point, all data must be written onto the virtual disk. Another program that mounts the file system at a later point in time must see the previously created files and the data that was written. This means that whenever *umount_fs* is called, all meta-information and file data (that you could temporarily have only in memory; depending on your implementation) must be written out to disk.

In addition to the management routines listed above, you are supposed to implement the following file system functions (which are very similar to the corresponding Linux file system operations). These file system functions require that a file system was previously mounted.

File system operations

<pre>int fs_open(char *name);</pre>	<p>Open file specified by <i>name</i> for reading and writing. Return -1 on failure, and a non-negative integer, which is a file descriptor that can be used to subsequently access this file when successful.</p> <p>Note: The same file (file with the same name in a directory) can be opened multiple times. When this happens, your file system is supposed to provide multiple, independent file descriptors. Your library must support a maximum of 32 file descriptors that can be open simultaneously. It is a failure when the file with <i>name</i> cannot be found (i.e., it has not been created previously or is already deleted). It is also a failure when there are already 32 file descriptors active. When a file is opened, the file offset (seek pointer) is set to 0 (the beginning of the file).</p>	<p style="text-align: center;">63' (Optional, each for 7')</p>
<pre>int fs_close(int fildes);</pre>	<p>Close file descriptor <i>fildes</i>. A closed file descriptor can no longer be used to access the corresponding file. Return 0 upon successful completion. In case the file descriptor <i>fildes</i> does not exist or is not open, the function returns -1.</p>	
<pre>int fs_create(char *name);</pre>	<p>Create a new file with name <i>name</i> in the root directory of your file system. The file is initially empty. The maximum length for a file name is 15 characters. Also, there can be at most 64 files in the directory. Return 0 on success, and -1 on failure. It is a failure when the</p>	

	<p>file with <i>name</i> already exists, when the file name is too long (it exceeds 15 characters), or when there are already 64 files present in the root directory.</p> <p>Note that to access a file that is created, it has to be subsequently opened.</p>	
<pre>int fs_delete(char *name);</pre>	<p>Delete the file with name <i>name</i> from the root directory of your file system and frees all data blocks and meta-information that correspond to that file. The file that is being deleted must not be open. That is, there cannot be any open file descriptor that refers to the file <i>name</i>. When the file is open at the time that <i>fs_delete</i> is called, the call fails and the file is not deleted.</p> <p>Return 0 upon successful completion, and -1 on failure. It is a failure when the file with <i>name</i> does not exist. It is also a failure when the file is currently open (i.e., there exists at least one open file descriptor that is associated with this file).</p>	
<pre>int fs_read(int fildes, void *buf, size_t nbyte);</pre>	<p>Attempt to read <i>nbyte</i> bytes of data from the file referenced by the descriptor <i>fildes</i> into the buffer pointed to by <i>buf</i>. The function assumes that the buffer <i>buf</i> is large enough to hold at least <i>nbyte</i> bytes. When the function attempts to read past the end of the file, it reads all bytes until the end of the file.</p> <p>Upon successful completion, the number of bytes that were actually read is returned. This number could be smaller than <i>nbyte</i> when attempting to read past the end of the file (when trying to read while the file pointer is at the end of the file, the function returns zero). In case of failure, the function returns -1. It is a failure when the file descriptor <i>fildes</i> is not valid.</p>	

	<p>The <i>read</i> function implicitly increments the file pointer by the number of bytes that were actually read.</p>	
<pre>int fs_write(int fildes, void *buf, size_t nbyte);</pre>	<p>Attempt to write <i>nbyte</i> bytes of data to the file referenced by the descriptor <i>fildes</i> from the buffer pointed to by <i>buf</i>. The function assumes that the buffer <i>buf</i> holds at least <i>nbyte</i> bytes. When the function attempts to write past the end of the file, the file is automatically extended to hold the additional bytes. It is possible that the disk runs out of space while performing a write operation. In this case, the function attempts to write as many bytes as possible (i.e., to fill up the entire space that is left). The maximum file size is 16M (which is, 4,096 blocks, each 4K).</p> <p>Upon successful completion, the number of bytes that were actually written is returned. This number could be smaller than <i>nbyte</i> when the disk runs out of space (when writing to a full disk, the function returns zero). In case of failure, the function returns -1. It is a failure when the file descriptor <i>fildes</i> is not valid. The <i>write</i> function implicitly increments the file pointer by the number of bytes that were actually written.</p>	
<pre>int fs_get_filesize(int fildes);</pre>	<p>Return the current size of the file pointed to by the file descriptor <i>fildes</i>. In case <i>fildes</i> is invalid, the function returns -1.</p>	
<pre>int fs_lseek(int fildes, off_t offset);</pre>	<p>Set the file pointer (the offset used for read and write operations) associated with the file descriptor <i>fildes</i> to the argument <i>offset</i>.</p> <p>It is an error to set the file pointer beyond the end of the file. To append to a file, one can set the file pointer to the end of a file, for example, by calling <i>fs_lseek(fd, fs_get_filesize(fd));</i>.</p>	

	Return 0 upon successful completion, and -1 on failure. It is a failure when the file descriptor <i>fildes</i> is invalid, when the requested <i>offset</i> is larger than the file size, or when <i>offset</i> is less than zero.	
int fs_truncate(int fildes, off_t length);	<p>Cause the file referenced by <i>fildes</i> to be truncated to <i>length</i> bytes in size. If the file was previously larger than this new size, the extra data is lost and the corresponding data blocks on disk (if any) must be freed. It is not possible to extend a file using <i>fs_truncate</i>. When the file pointer is larger than the new <i>length</i>, then it is also set to <i>length</i> (the end of the file).</p> <p>Return 0 upon successful completion, and -1 on failure. It is a failure when the file descriptor <i>fildes</i> is invalid or the requested <i>length</i> is larger than the file size.</p>	

iii. Implementation

In principle, you can implement the file system in any way that you want (as long as 4,096 blocks of the disk remain available to store file data). However, it might be easier when you borrow ideas from existing file system designs. You can model your file system after the FAT (file allocation table) design, although it is also possible to use a Unix (inode)-based design.

In general, you will likely need a number of data structures on disk, including a super block, a root directory, information about free and empty blocks on disk, file meta-information (such as file size), and a mapping from files to data blocks.

The *super block* is typically the first block of the disk, and it stores information about the location of the other data structures. For example, you can store in the super block the whereabouts of the file allocation table, the directory, and the start of the data blocks.

The *directory* holds the names of the files. When using a FAT-based design, the directory also stores, for each file, its file size and the head of the list of corresponding data blocks. When you use inodes, the directory only stores the mapping from file names to inodes.

The *file allocation table (FAT)* can be used to keep track of empty blocks *and* the mapping between files and their data blocks. When you use an inode-based design, you will need a bitmap to mark disk blocks as used and an inode array to hold file information (including the file size and pointers to data blocks).

In addition to the file-system-related data structures on disk, you also need support for file descriptors. A file descriptor is an integer in the range between 0 and 31 (inclusive) that is returned when a file is opened, and it is used for subsequent file operations (such as reading and writing). A file descriptor is associated with a file, and it also contains a file offset (seek pointer). This offset indicates the point in the file where read and write operations start. It is implicitly updated (incremented) whenever you perform a *fs_read* or *fs_write* operation, and it can be explicitly moved within the file by calling *fs_lseek*. Note that file descriptors are not stored on disk. They are only meaningful while an application is running and the file system is mounted. Once the file system is unmounted, file descriptors are no longer meaningful (and, hence, should be all closed before a call to *umount_fs*).

iv. Due Dates and Deliverables

Please follow the instructions below exactly!

1. The project is an individual project. It is due on November 27, 2016 at 5pm.
2. Your files must be in a directory whose naming format is xxx(student ID)-xxx(name).zip(rar, tar...).

3. The implementation must be done in C/C++. Please do not include the code from *disk.c* in your library. We will link our *disk.o* against your code and our test applications. Of course, you can (and should) include *disk.h*. Moreover, you can (and should) use *disk.c* for local testing.
4. All files that you need to build your library must be included (sources, headers, makefile) in that folder. We will just call make and expect that the object file *fs.o* is built from your sources. We will test and link must be called *fs.o*.
5. (Score 7') Include a README with this project. Explain how you thought and what you did briefly. All sources and headers files should have code comments necessary.
6. All deliverables should be submitted via the department ftp.
7. Any cheat behaviors will get zero.